

Feedback Directed Helper Threaded Inter-Core Data Prefetching

Min Cai, Zhimin Gu

Abstract—Helper threaded inter-core data prefetching on shared cache chip multiprocessors (CMPs), in its most basic form, (1) uses the shared cache enabled fast inter-thread communication when both the main core and the otherwise idle neighboring helper core access the same cache line; (2) constructs a helper thread in the software running on the helper core to speculatively issue LLC prefetch requests to the predicted memory addresses before the computation core access these addresses. Effective helper threaded inter-core data prefetching demands that the helper thread (HT) should issue correct and timely LLC requests just before the main thread (MT) requests them. Unfortunately, this ideal case can not be assumed in the existing implementation of helper threaded inter-core data prefetching: inaccurate and/or untimely LLC requests coming from HT could not contribute to the MT performance, but instead stress and pollute LLC if no effective LLC replacement and pollution-aware feedback techniques are employed.

In this paper, we present a pollution-aware feedback mechanism for dynamic helper threaded inter-core data prefetching, based on a pollution-aware taxonomy of HT LLC requests. A pollution aware taxonomy of HT LLC requests is presented from the view of the contribution of HT LLC requests to the MT performance. Based on the pollution aware taxonomy of HT LLC requests, the feedback mechanism for dynamic helper threaded inter-core data prefetching is proposed to finetune LLC replacement and parameters of the HT scheme. Experimental results from cycle accurate simulation of the Pthreads based helper threaded version of the memory-intensive mst benchmark in Olden show that: (1) there is non-trivial LLC pollution caused by HT in helper-threaded data prefetching on CMPs; (2) LLC request taxonomy based dynamic LLC replacement can improve the effectiveness and timeliness of helper-threaded data prefetching on CMPs.

Index Terms—chip multiprocessors, helper threaded inter-core data prefetching, cache replacement, cache pollution, inter-thread interference

I. INTRODUCTION

HELPER threaded inter-core data prefetching on shared cache chip multiprocessors (CMPs), in its most basic form, (1) uses the shared cache enabled fast inter-thread communication when both the main core and the otherwise idle neighboring helper core access the same cache line; (2) constructs a helper thread in the software running on the helper core to speculatively issue LLC prefetch requests to the predicted memory addresses before the computation core access these addresses. Effective helper threaded data prefetching demands that the helper thread (HT) should issue correct and timely LLC requests just before the main thread (MT) requests them. Unfortunately, this ideal case can not

be assumed in the existing implementation of helper threaded inter-core data prefetching: inaccurate and/or untimely LLC requests coming from HT could not contribute to the MT performance, but instead stress and pollute LLC if no effective LLC replacement techniques are employed.

Several metrics have been proposed in the past for evaluating the effectiveness of hardware based data prefetching, among which prefetch accuracy and coverage are the most intuitive ones [1]. We can easily adapt the definitions of accuracy and coverage for hardware based data prefetching to the HT scheme. HT request accuracy is defined as the ratio of the number of useful HT requests to the number of total HT requests. And HT request coverage is defined as the ratio of the number of useful HT requests to the number of MT misses plus MT hits to HT requested data. Here, an HT request is called useful when its requested data is referenced by MT before evicted. Furthermore, similar to the approach to hardware prefetching, to measure coverage and accuracy, all HT requests can be categorized into “useful” and “useless” HT requests [1]. A “useful” HT request is one whose brought data is hit by a MT request before it is replaced, while a “useless” HT request is one whose brought data is replaced before it is hit by a MT request. However the above traditional accuracy and coverage metrics for HT requests and the classification of “useful” and “useless” HT requests don’t care about the LLC pollution and inter-thread interference caused by HT LLC requests, which are two kinds of deficiencies in the HT scheme.

In this paper, we present a pollution-aware feedback mechanism for dynamic helper threaded inter-core data prefetching, based on a pollution-aware taxonomy of HT LLC requests. A pollution aware taxonomy of HT LLC requests is presented from the view of the contribution of HT LLC requests to the MT performance. Based on the pollution aware taxonomy of HT LLC requests, the feedback mechanism for dynamic helper threaded inter-core data prefetching is proposed to finetune LLC replacement and parameters of the HT scheme. Experimental results from cycle accurate simulation of the Pthreads based helper threaded version of the memory-intensive mst benchmark in Olden show that: (1) there is non-trivial LLC pollution caused by HT in helper-threaded data prefetching on CMPs; (2) LLC request taxonomy based dynamic LLC replacement can improve the effectiveness and timeliness of helper-threaded data prefetching on CMPs. We assume here a two level cache hierarchy where L1 caches are private and the L2 cache is shared among all processor cores on a single chip.

The main contributions of this paper can be summarized as answers for the following two questions:

Min Cai, School of Computer Science and Technology, Beijing Institute of Technology, Beijing, China, e-mail: min.cai.china@gmail.com.

Zhimin Gu, School of Computer Science and Technology, Beijing Institute of Technology, Beijing, China, e-mail: zmggu@x263.net.

- 1) How can the HT LLC requests be classified based on the cache pollution and inter-thread LLC interference caused by HT LLC requests in helper threaded inter-core data prefetching (or the HT scheme)?
- 2) How can LLC replacement be improved based on the observed result of the proposed pollution-aware HT LLC request taxonomy in the HT scheme?

The rest of this paper is organized as follows. Section 2 presents the pollution aware taxonomy of HT requests and the experimental result of mst in Olden. Section 3 discusses the dynamic LLC replacement for helper threaded inter-core data prefetching and its experimental results. Section 4 presents the experimental methodology used in this work. Section 5 discusses results. Section 6 talks about related work. In section 7 we conclude the paper.

II. POLLUTION AWARE TAXONOMY OF HT LLC REQUESTS

A. The Notion of Inter-Thread Reuse Distance in the HT Scheme

The notion of reuse distance can shed light on how a pollution-aware taxonomy of HT LLC requests can be constructed. The traditional intra-thread reuse distance of a reference to data element x is defined as the number of distinct data elements that have been referenced between two contiguous accesses of x (or ∞ if the element has not been referenced thereafter). Thus, these two contiguous accesses are issued by the same thread in the traditional notion. To accomodate the case where one thread T_b accesses data element x that has been previously brought into the cache by another thread T_a , we introduce the notion of T_a - T_b inter-thread reuse distance, as compared to the traditional intra-thread reuse distance applied to either T_a or T_b .

To show why the notion of inter-thread reuse distance is important in the HT scheme as compared to the traditional notion of intra-thread reuse distance, we can consider the mst benchmark in the Olden suite. Its pointer traversing code structure inherently exhibits irregular memory access pattern in its original, single threaded version, which renders the common LRU replacement policy inefficient to reduce LLC misses. Here we use RD_MT to refer to the intra-thread reuse distance in the main thread, and $ITRD_HT_MT$ to refer to the inter-thread reuse distance between HT and MT where a data element is first brought to LLC by HT, and later on used by MT. Therefore, the values of RD_MT in most MT LLC requests is high which reflects the irregular memory access pattern in mst. However, the values of $ITRD_HT_MT$ in HT LLC requests should be very low when the HT scheme is efficient to reduce LLC misses in MT; otherwise, high values of $ITRD_HT_MT$ indicate the inefficiency of the HT scheme where most data brought by HT is replaced before used by MT.

B. Pollution Aware Taxonomy of HT LLC Requests

Based on the notions of HT-MT inter-thread reuse distance (HT_MT_ITRD) and MT intra-thread reuse distance

(MT_RD), we can construct a pollution aware taxonomy of HT LLC requests where HT LLC requests are classified into three types: good, bad and ugly.

Let's assume that when an HT LLC request with data d evicts the LLC line containing victim data v which is brought by MT, and afterwards the LLC line containing d is evicted by an MT LLC request with data m , then we can see that:

- 1) $HT_MT_ITRD(d) < MT_RD(v) < MT_RD(m)$, indicating it's a *good HT request* whose data is hit by MT before evicted. It has positive impact on MT performance since it reduces one MT miss;
- 2) $MT_RD(v) < HT_MT_ITRD(d) < MT_RD(m)$, indicating it's a *bad HT request* whose data displaces an LLC line that will later be needed by MT. It is harmful for MT performance which should be prevented as much as possible;
- 3) $MT_RD(m) < HT_MT_ITRD(d)$, $MT_RD(m)$, indicating it's an *ugly HT request* which is not either good or bad. It has little performance impact on MT performance because the requested data is not referenced by MT before evicted and it does not evict any data that will be used by MT.

Good HT LLC requests can be further divided into timely requests and late requests. An HT request is called *late HT request* if the data is hit in LLC MSHRs by an MT request where it is brought from memory by HT. Late HT requests only partially hide the LLC miss latency in MT requests, but they are good indicators for potential performance improvement of the HT scheme because late HT requests may be converted to timely or normal *good HT requests* by finetuning the parameters of the HT scheme.

C. Hardware Support for the Pollution-Aware Taxonomy of HT LLC Requests

1) *LLC request and replacement event tracking*: To monitor the request and replacement activities in LLC, we need to consider the event when the LLC receives a request coming from the upper level cache, whether it is a hit or a miss. The event has a few important properties to be used in the experiment, e.g., the address of the requested LLC line, the requester memory hierarchy access, line found in the LLC, a boolean value indicating whether the request hits in the LLC, and a boolean value indicating whether the request needs to evict some LLC line. This event is similar to one used in [2].

2) *LLC HT request state tracking*: In order to track the HT request states in the LLC, we need to add one field to each LLC line to indicate whether the line is brought by the main thread (MT) or the helper thread (HT) or otherwise invalid (INVALID).

3) *LLC HT request victim state tracking*: In order to track victims replaced by HT requests, we need to add an LRU cache named HT Request Victim Cache (HTRVC) to maintain the LLC lines that are evicted by HT requests. Similar to the evict table used in [2], the HTRVC has the same structure of the LLC, but there is no direct mapping between LLC lines and HTRVC lines. HTRVC has only the purpose of profiling, so it has no impact on performance.

4) *Detecting Late HT Requests:* Lastly, in order to measure late HT requests, we only need to identify the event when an MT request hits to an LLC line which is being brought by an inflight HT request coming from the upper level cache. This can be accomplished by monitoring the LLC MSHRs or equivalent hardware components.

D. Algorithms for Tracking HT LLC Requests and Victims

There are two invariants that should be maintained:

1) # of HT Lines in the LLC Set = # of Victim Entries in the HTRVC Set;

a) # of Victim Entries in HTRVC Set + # of Valid MT LLC Lines in Set \leq LLC Set Associativity.

HT lines refer to the cache lines that are brought by HT requests. From the above two invariants, we can easily conclude that: # of HT Lines in the LLC Set + # of Valid MT LLC Lines in Set \leq LLC Set Associativity. Actions should be taken in LLC and HTRVC when filling an LLC line or servicing an incoming LLC request.

1) *Actions taken on Filling an LLC Line:* When filling an LLC line, we should consider five cases:

- An HT request evicts an INVALID line. In this case, no eviction is needed.
- An HT request evicts an LLC line which is previously brought by an MT request. In this case, eviction is needed to make room for the incoming HT request.
- An HT request evicts an LLC line which is previously brought by an HT request. In this case, eviction is needed to make room for the incoming HT request.
- An MT request evicts an LLC line which is previously brought by an HT request. In this case, eviction is needed to make room for the incoming MT request.
- An MT request evicts an LLC line which is previously brought by an MT request, and there exists one line that is brought by an HT request. In this case, eviction is needed to make room for the incoming MT request.

Specific actions taken on the above five cases are listed in Fig.1, where `hitInLLC`: whether the request hits in LLC or not, `requesterIsHT`: whether the request comes from HT or not, `hasEviction`: whether the request needs to evict some data, `lineFoundIsHT`: whether the LLC line found is brought by HT or not, and `htRequestFound()`: whether there is at least one line in the LLC set that is brought by HT.

2) *Actions taken on Servicing an Incoming LLC Request:* When servicing an incoming LLC request, either hit or miss, we should consider four cases:

- LLC miss and victim hit, which indicates a bad HT request. This happens when HT request evicts useful data.
- HT LLC hit, which indicates a good HT request. This happens when HT requested data is hit by MT request before evicted data.

```
//HT miss
if(!hitInLLC && requesterIsHT) {
    totalHtRequests++;
}

//Case 1
if(requesterIsHT && !hitInLLC && !hasEviction) {
    llc.setHT(set, llcLine.way);
    htrvc.insertNullEntry(set);
}
//Case 2
else if(requesterIsHT && !hitInLLC && hasEviction && !
    lineFoundIsHT) {
    llc.setHT(set, llcLine.way);
    htrvc.insertDataEntry(set, llcLine.tag);
}
//Case 3
else if(requesterIsHT && !hitInLLC && hasEviction &&
    lineFoundIsHT) {
}
//Case 4
else if(!requesterIsHT && !hitInLLC && hasEviction &&
    lineFoundIsHT) {
    llc.setMT(set, llcLine.way);
    htrvc.removeLRU(set);
}
//Case 5
else if(!requesterIsHT && !lineFoundIsHT) {
    if(htRequestFound()) {
        htrvc.removeLRU(set);
        htrvc.insertDataEntry(set, llcLine.tag);
    }
}
```

Figure 1. Actions Taken When Filling an LLC Line

- HT LLC hit and victim hit. This happens when useful data is evicted and brought back in by HT request.
- MT LLC hit and victim hit. This happens when useful data is evicted and brought back in by HT request and hit to by MT request.

Specific actions taken on the above five cases are listed in Fig.2, where `mtHit`: whether the request comes from MT and hits in the LLC, `htHit`: whether the request comes from HT and hits in the LLC, and `vtHit`: whether the request comes from MT and hits in the HTRVC.

```
//Case 1
if(!mtHit && !htHit && vtHit) {
    badHtRequests++;
    htrvc.setLRU(set, vtLine.way);
}
//Case 2
else if(!mtHit && htHit && !vtHit) {
    llc.setMT(set, llcLine.way);
    goodHtRequests++;
    htrvc.removeLRU(set);
}
//Case 3
else if(!mtHit && htHit && vtHit) {
    llc.setMT(set, llcLine.way);
    htrvc.setLRU(set, vtLine.way);
    htrvc.removeLRU(set);
}
//Case 4
else if(mtHit && !htHit && vtHit) {
    htrvc.setLRU(set, vtLine.way);
}
```

Figure 2. Actions Taken When Servicing an LLC Request

III. POLLUTION-AWARE FEEDBACK DIRECTED HELPER THREADED INTER-CORE DATA PREFETCHING

Based on the pollution-aware taxonomy of HT LLC requests, we propose a dynamic XXX..

IV. EVALUATION METHODOLOGY

We use an in-house CMP architectural simulator named Archimulator in our experiments mentioned in this work. Archimulator is a flexible execution-driven architectural simulator written in Java and running on Linux. It provides fast forward functional simulation and cycle-accurate application-only simulation of MIPSII executables on multicore architectures consisting of out-of-order super-scalar cores and configurable memory hierarchy of directory-based MESI coherence. It has basic support of simulating Pthreads based multithreaded workloads.

A. Simulated CMP Architecture

As shown in Fig.3, the simulated target CMP architecture has two cores where each core is a two-way SMT with its own private L1 caches (32KB 8-way data caches and 32KB 4-way instruction caches). Both cores share a 4MB 8-way L2 cache. MESI coherence is maintained between L1 caches. An LRU cache called HTRVC is attached to the LLC(L2) to implement the taxonomy of HT LLC requests, which will be detailed in the next section. Detailed microarchitecture parameters are listed in Tab.I.




Figure 3. Simulated CMP Architecture

B. Software Context to Hardware Thread Mapping

In a typical Pthreads based HT program, there are three threads when running: MT, HT and the Pthreads manager thread. The Pthreads manager thread takes the role of spawning, suspending and resuming HT by passing signals to HT. Consider a simulated target multicore machine which has two cores where each core supports two hardware threads. In our application-only simulation using Archimulator, without the

Pipeline	4-wide superscalar OoO 2 x 2 CMP; physical register file capacity: 128; decode buffer capacity: 96; reorder buffer capacity: 96; load store queue capacity: 48				
Branch Predictors	Perfect branch predictor				
Execution Units	Name	Count	Operation Lat.	Issue Lat.	
	Int ALU	8	2	1	
	Int Mult	2	3	1	
	Int Div		20	19	
	Fp Add	8	4	1	
	Fp Compare		4	1	
	Fp Convert		4	1	
	Fp Mult	2	8	1	
	Fp Div		40	20	
	Fp Sqrt		80	40	
	Read Port	4	1	1	
	Write Port		1	1	
Cache Geometries	Name	Size	Assoc.	Line Size	Hit Lat.
	l1i	32KB	4	64B	1
	l1d	32KB	8	64B	1
	l2	4096KB	8	64B	10
Interconnect	Switch based P2P topology, 32B link width				
Main Memory	4GB, 200-cycle fixed latency				

Table I
BASELINE HARDWARE CONFIGURATIONS

OS intervention, one hardware thread can only run at least one software context (or simply called thread). Therefore, the typical software context to hardware thread mappings can be: C0T0 → MT, C0T1 → Pthreads manager thread, C1T0 → HT (C = core, T = thread), as shown in Fig.3. We use this context mapping in the following discussions.

C. ROI Based Two-Phase Fast Simulation of Helper Threaded Workloads

As shown in Fig.4, here we use two-phase simulation strategy to reduce simulation time of memory-intensive benchmark executions which often take very long to complete. We utilize the special case when a MIPS32 assembly instruction `addiu` is invoked with register operands being R0: `addiu R0, R0, imm` to indicate the spawning point of the helper thread (named by “HT Spawn”: `addiu R0, R0, 3720`). Since the code executed after encountering the “HT spawn” is our region of interest (ROI) in this work, code execution before encountering the “HT spawn” can be simulated functionally without the hassle of modeling the microarchitecture details while not hurting the experimental results. To ensure caches are warmed up, we simulate in the detailed simulation mode 100 million instructions in the main thread after the “HT Spawn” pseudocall is encountered.

V. RESULTS

We use the original version and helper threaded version of mst as workload used in our experiments. All versions of mst benchmark are cross-compiled with gcc flag “-O3”. Input args are “4000”.

Tab.II summarizes overall performance of the benchmark in which all the speedups are computed against the baseline unmodified version of the mst; K = lookahead, P = blocksize.

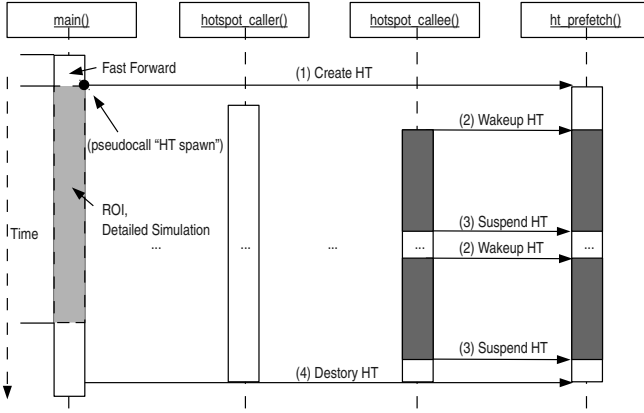


Figure 4. ROI Based Two Phase Fast Simulation of Helper Threaded Inter-Core Data Prefetching

VI. RELATED WORK

Our taxonomy of HT LLC requests is mostly similar to the work in [2], which presents a taxonomy of hardware prefetches based on the idea of shared cache pollution in the hardware based data prefetching for shared L2 CMP. A hardware structure called the Evict Table (ET) is attached to the LLC to gauge the amount of shared cache pollution caused by hardware prefetching. Good, bad and ugly requests are identified based on cache replacement activities involved by hardware prefetches.

Our feedback-directed scheme is similar to the work in [3], which proposes a feedback directed mechanism for hardware prefetching.

VII. CONCLUSION

TODOs: our work: what? how? result? Further work?

ACKNOWLEDGMENTS

This work was supported by the National Natural Science Foundation of China under the contract No. 61070029.

REFERENCES

- [1] V. Srinivasan, E. S. Davidson, and G. S. Tyson, "A prefetch taxonomy," *IEEE Trans. Computers*, vol. 53, no. 2, pp. 126–140, 2004.
- [2] B. Mehta, D. Vantrease, and L. Yen, "Cache showdown: The good, bad and ugly," Tech. Rep., 2004.
- [3] S. Srinath, O. Mutlu, H. Kim, and Y. N. Patt, "Feedback directed prefetching: Improving the performance and bandwidth-efficiency of hardware prefetchers," in *Proc. 13th International Conference on High-Performance Computer Architecture (13th HPCA'07)*. San Francisco, CA, USA: IEEE Computer Society, feb 2007, pp. 63–74.

Table II
MST HT LLC REQUEST TAXONOMY RESULTS UNDER LRU LLC REPLACEMENT

K, P	cycles	speedup	MT LLC Read Misses	HT LLC Requests			
				Good	Bad	Ugly	Late
L2: 1M, 4-way							
Unmodified	2128177521						
10, 10	1413640215			5314478	9	15831	2291709
20, 10	1201000213			4596361	25	33536	571166
20, 20	1391083752			5132796	0	514	2007020
320, 320	1299136860			4629522	449	95403	1049485
640, 320	1281837899			3658166	1252	720821	190
640, 640	1342389544			4021349	1697	1053477	633125
L2: 1M, 8-way							
Unmodified	2128379206						
10, 10	1413844837			5300159	35	28713	2278392
20, 10	1201909631			4602015	28	28613	580520
20, 20	1394465307			5110279	86	32673	1996862
320, 320	1298792424			4642918	577	62818	1059197
640, 320	1281744764			3661219	1939	703163	226
640, 640	1337983575			4043866	2526	1017277	633273
L2: 2M, 4-way							
Unmodified	2120356167						
10, 10	1399634774			5257629	26	30005	2305058
20, 10	1174856777			4516140	0	319	488142
20, 20	1373584291			5031697	16	11407	1943799
320, 320	1309023530			4724771	83	9234	1302455
640, 320	1310880244			3462836	595	670785	35252
640, 640	1336338360			4244832	714	681175	932923
L2: 2M, 8-way							
Unmodified	2120668282						
10, 10	1394816682			5280283	0	295	2335138
20, 10	1167492039			4485981	0	522	446726
20, 20	1374797661			5039645	6	8408	1972291
320, 320	1321725531			4753262	130	597	1410322
640, 320	1293183359			3862339	302	77499	388081
640, 640	1315342952			4673206	354	75351	1302622
L2: 4M, 4-way							
Unmodified	2114816878						
10, 10	1389796215			5250168	0	131	2323420
20, 10	1142128970			4396365	0	167	274300
20, 20	1368304714			5003078	0	178	1945235
320, 320	1325248401			4720657	34	679	1434343
640, 320	1319660419			3510111	403	499676	183632
640, 640	1335035605			4334291	356	504965	1078711
L2: 4M, 8-way							
Unmodified	2112469520						
10, 10	1387072573			5235217	4	5389	2320005
20, 10	1159240286			4451922	0	284	435001
20, 20	1369694579			4991352	2	6000	1958224
320, 320	1330619790			4720528	21	458	1478493
640, 320	1293317577			3884985	161	11982	464030
640, 640	1319089645			4692382	195	10284	1393115

Table III
MST HT LLC REQUEST TAXONOMY RESULTS UNDER DYNAMIC LLC
REPLACEMENT

K, P	cycles	speedup	MT LLC Read Misses	HT LLC Requests			
				Good	Bad	Ugly	Late
L2: 1M, 4-way							
Unmodified							
10, 10	1415822623			5289283	24	49899	2277062
20, 10	1198389681			4602585	14	17316	564991
20, 20	1391086968			5133595	0	516	2007792
320, 320	1297462890			4643525	455	78974	1055340
640, 320	1281823292			3658069	1279	721139	169
640, 640	1342388863			4020819	1708	1053730	632577
L2: 1M, 8-way							
Unmodified							
10, 10	1414936009			5287475	46	45585	2270900
20, 10	1201422011			4610353	23	17889	586688
20, 20	1394761071			5104755	85	36819	1992954
320, 320	1298818263			4642534	560	65117	1059315
640, 320	1281878608			3659494	1953	708409	249
640, 640	1339670226			4037280	2490	1030596	634789
L2: 2M, 4-way							
Unmodified							
10, 10	1399638510			5257106	25	30014	2304490
20, 10	1177357943			4522885	0	317	507372
20, 20	1373593890			5033452	24	11405	1945536
320, 320	1309033612			4724233	81	9391	1301979
640, 320	1310831837			3462866	586	671131	35181
640, 640	1336354341			4243643	683	681561	931757
L2: 2M, 8-way							
Unmodified							
10, 10	1394822639			5280466	0	292	2335433
20, 10	1168890592			4489385	0	538	456889
20, 20	1374802543			5039894	6	8407	1972496
320, 320	1321707262			4753309	128	622	1410366
640, 320	1293146652			3862072	288	77652	387710
640, 640	1315353076			4674418	378	75388	1303841
L2: 4M, 4-way							
Unmodified							
10, 10	1389796412			5250450	0	132	2323455
20, 10	1140630504			4391249	0	163	261977
20, 20	1368310342			5001750	0	177	1943891
320, 320	1325261372			4720049	33	804	1433818
640, 320	1319659301			3509582	412	499934	183136
640, 640	1335042038			4337037	361	505222	1081397
L2: 4M, 8-way							
Unmodified							
10, 10	1387074688			5234229	4	5388	2319043
20, 10	1158143928			4447738	0	276	425326
20, 20	1369698439			4991788	3	6000	1958674
320, 320	1330618620			4720103	21	458	1478065
640, 320	1293325892			3886050	162	12020	465096
640, 640	1319104804			4692545	196	10290	1393320