

A Comparison of Two Approaches to Parallel Simulation of Multiprocessors

Andrew Over*, Bill Clarke, and Peter Strazdins

CC-NUMA Project,
Department of Computer Science,
The Australian National University

ISPASS 2007, 26 April 2007

<http://alto.anu.edu.au/~andrew/seminars/ispass.pdf>

1 Overview

- Motivation
- Simulator background and design
- Parallelization approaches:
 - via locking
 - via PDES techniques
- Performance comparison:
 - Methodology and workloads
 - Simulator metric stability
 - Performance comparison
 - Load balancing concerns
- Conclusions


2 Motivation

- CMP systems are now widely available:
 - Multi-core simulation is now needed in more situations
 - Common platforms are suitable hosts for parallel simulation
- Serial simulation of MP or CMP systems incurs a linear slowdown
- Obtaining results quickly is sometimes more useful than obtaining results efficiently
- Difficult to parallelize owing to degree of communication between processors on shared-memory platforms

3 Simulator Infrastructure

- Models multiprocessor UltraSPARC III Cu systems
 - In-order superscalar processor
 - 32KB I-cache, 64KB 4-way pseudo-random D-cache
 - 1–8MB 2-way E-cache
- Execution driven (fetch-decode-execute loop)
- Memory system model (including FirePlane interconnect protocol)
- Runs unmodified Solaris binaries using a syscall emulation layer
- Typical slowdown 500–1000×
- Simulation target Sun V1280:
 - 12 × 900MHz UltraSPARC III Cu
 - 24GB RAM, single snooping coherence domain
- Simulation target doubles as Simulation host

4 Simulator Implementation Details

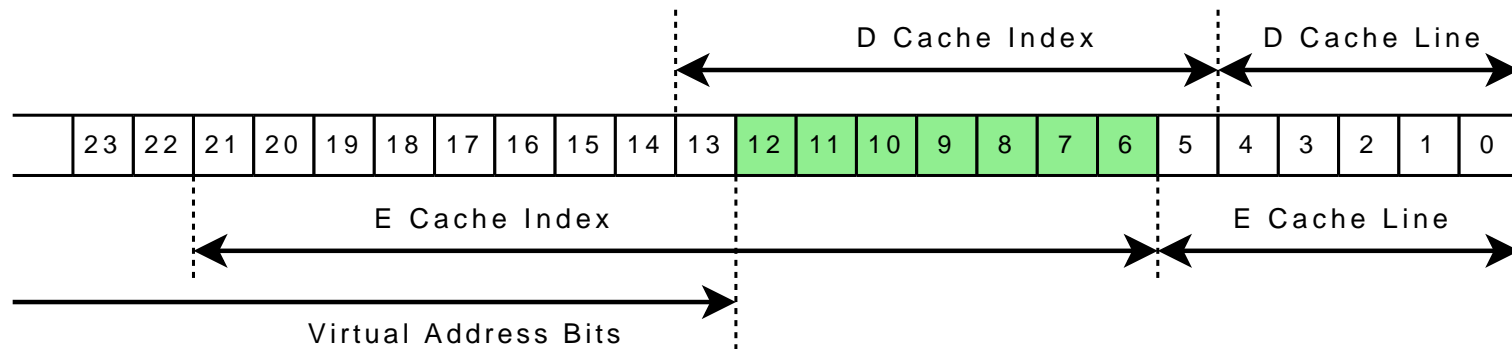


- Memory hierarchy is the key obstacle to parallelization
 - All processor interactions occur through cache coherence
 - Coherency protocol must operate in parallel but preserve correctness
- Processor model isolated from memory hierarchy by a *bridge*, which interacts with the *backplane* model
- Two different interconnect models:
 - Simple approximate model (*passive*)
 - Detailed complex model (*active*)

5 Simulator Passive Backplane

- Passive backplane is a simple interconnect model:
 - Fixed request latency based on type and coherency result
 - Coherency managed by directly inspecting and modifying foreign caches
 - No queueing of cache events (handled immediately)
- Serial multiprocessor simulation employs round-robin timeslicing
- Multiple simultaneous accesses \Rightarrow inconsistent state
- Host locking used to protect data structures
- Anecdotal 20% of instructions reference memory \Rightarrow a single lock would cripple performance

6 Simulator Passive Backplane (cont)



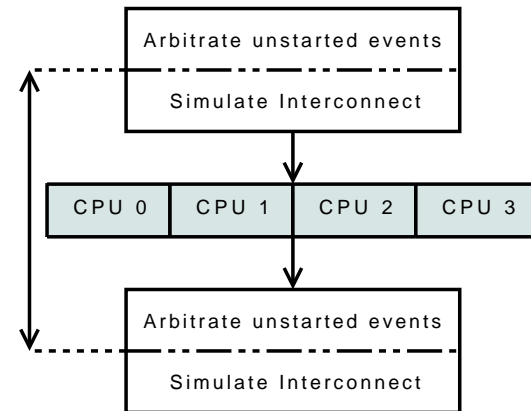
- Lock chosen according to cache index to prevent conflict
- No lock required on D-cache load hit
- Lock required on D-cache store, E-cache access
- Lock selection restricted to PA bits in D, E-cache indexes (128 locks)
- Majority of memory references remain unlocked
- Host locking introduces non-determinism
- Strict temporal ordering of requests not guaranteed

7 Simulator Active Backplane

- Detailed event-driven model of the interconnect:
 - All coherency states modelled (including transient states) based on FirePlane protocol
 - Cache coherency implemented via snoop request and snoop response events
- A minimum latency exists between event initiation and visibility on foreign processors (7 bus cycles @ 150MHz)
- Conservative parallel discrete event simulation techniques (PDES) may be applied
- Processors may run for limited periods without communication, before exchanging events

8 Simulator Active Backplane (cont)

- Parallel simulation consists of alternating serial and parallel phases:
 - Backplane runs in isolation
 - Processor models run in parallel interacting only with backplane

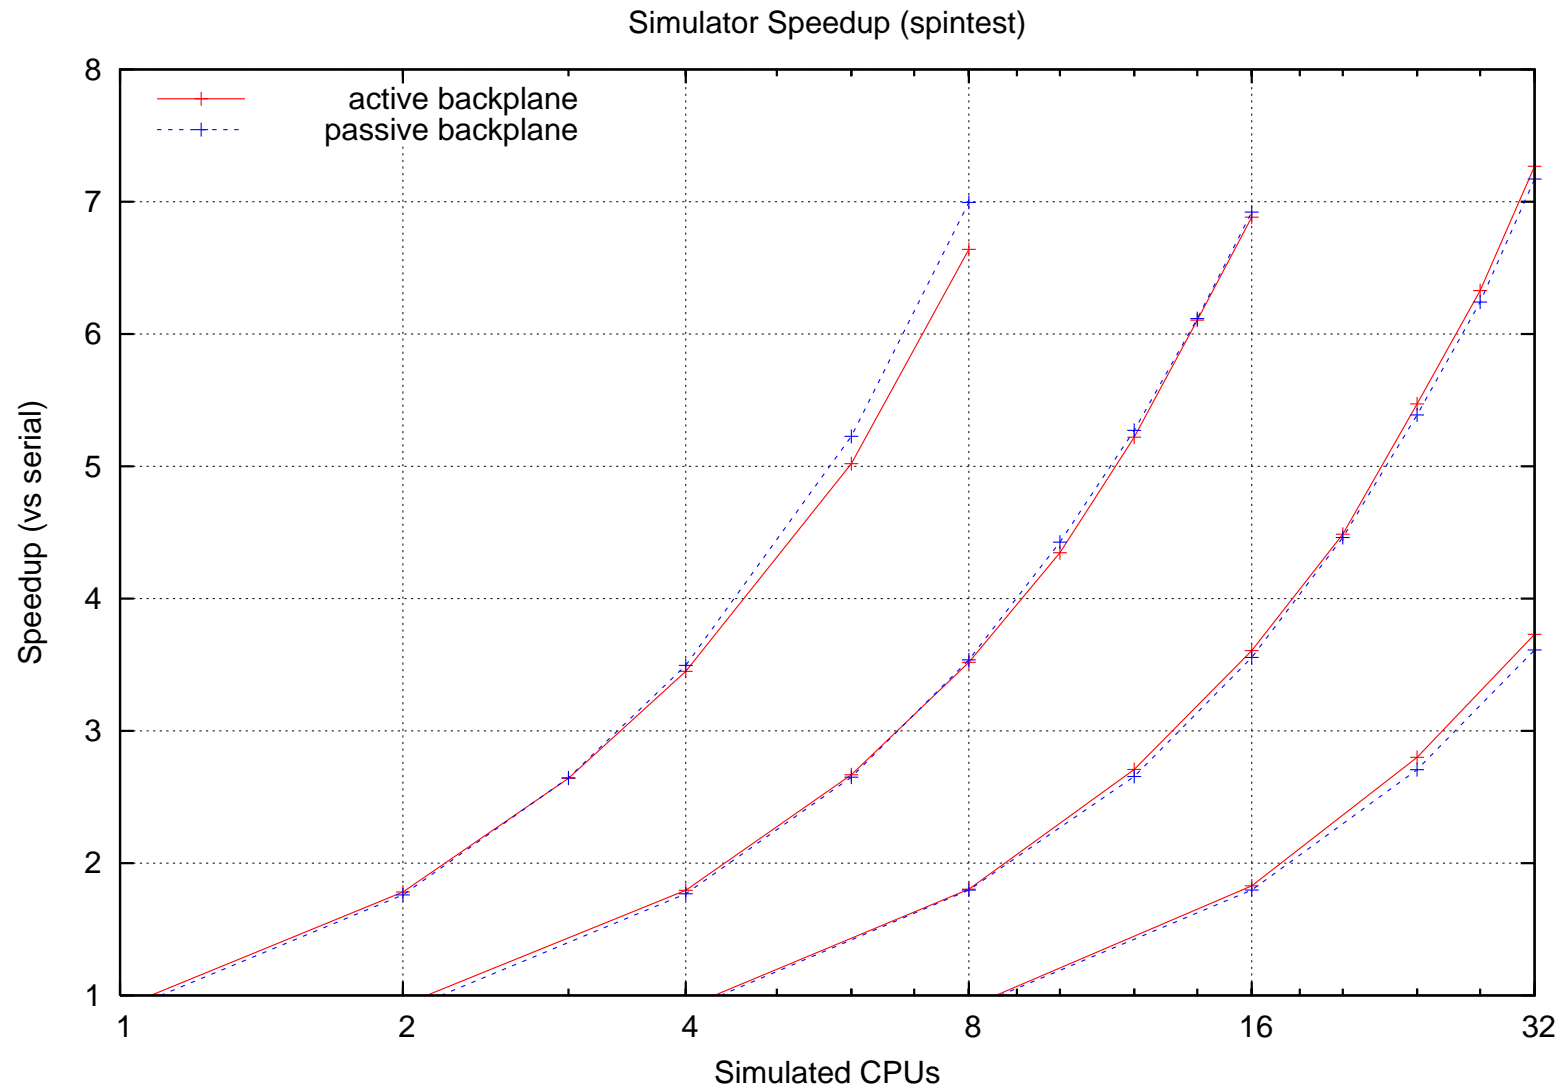


- Backplane arbitrates and queues events for processors
- Processor model handles queued events, forward responses
- All communication handled during serial phase
- Parallel processor simulation with complete determinism
- User-space barriers employed to synchronize at the end of a timeslice

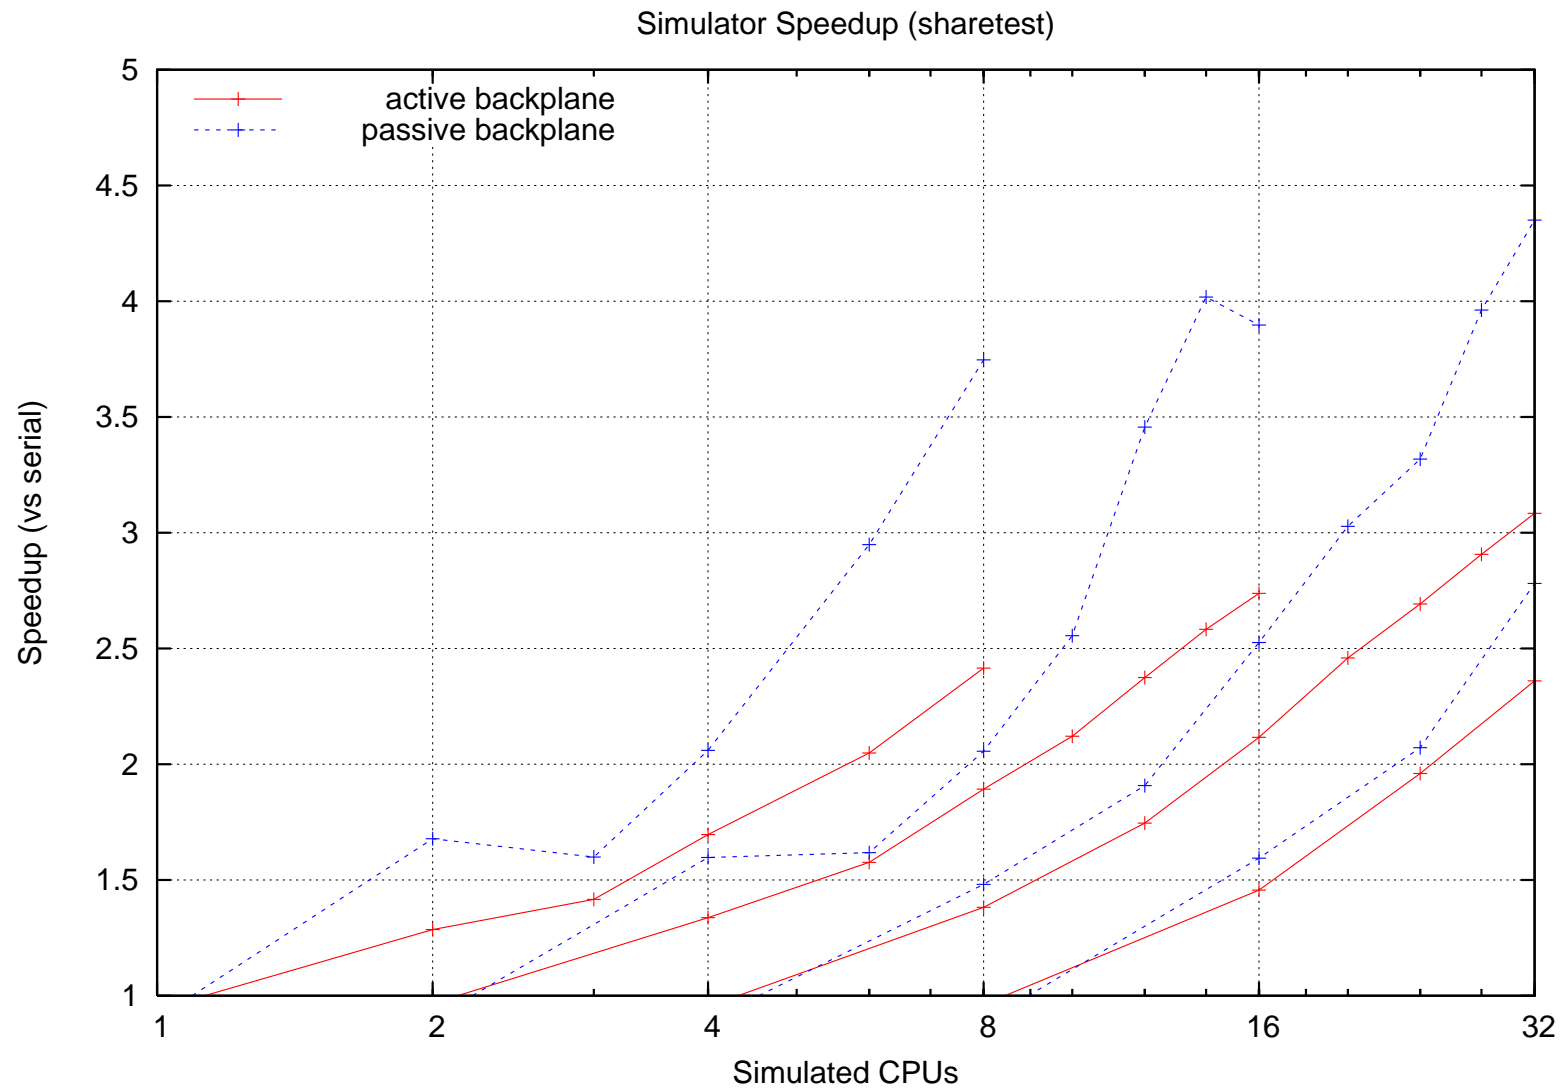
9 Methodology


- Parallel simulation of workloads normalized against serial simulation of an identical workload
- Two synthetic benchmarks used:
 - *spintest* – Tight loop, no memory references
 - *sharetest* – Frequent stores to small shared region
- Selected small scale NAS Parallel Benchmarks used
- Evaluating simulator performance, not simulation target performance
- Many simulation configurations tested (differing target processors per thread)

10 *spintest*




11 *sharetest*




12 *is.W*

13 *Passive Backplane Stability*



15 *Passive Backplane Stability*



16 *Conclusions*

- (Small scale) parallelization of architectural simulation is possible
- Non-deterministic parallelism conceptually simple with slight dependence on target architecture
- Deterministic parallelism is more involved and exploits interconnect properties
- Non-determinism introduces concerns over result stability
- Observed speedups depend on workloads and constrained by load-balancing of simulated processors

17 Future Work and Acknowledgements

- Future Work:
 - Additional analysis of stability concerns (including serial round-robin interleaving)
 - Evaluation of larger workloads, different architectures
- Acknowledgements:
 - CC-NUMA Project
 - Australian Research Council
 - Sun Microsystems
 - Gaussian Inc
 - ISPASS